



(a) Before: uniform, 64 λ per link

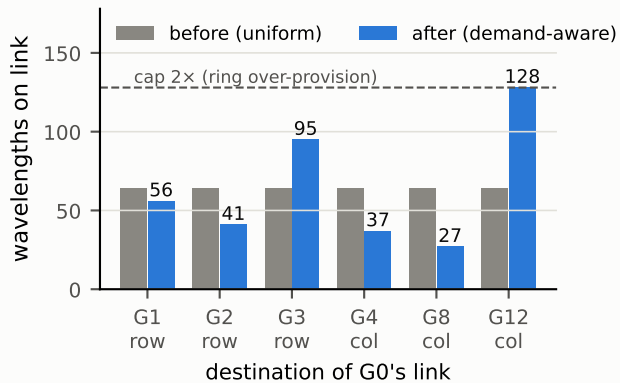


(b) After: demand-aware (floor 0.1, cap 2 $\times$)



orange = hot-expert GPUs (load $\geq 1.5\times$); blue = source G0; arrow width $\propto \lambda$

(c) Source G0: λ per outgoing link



(d) After: active drop rings per GPU

