



(a) Before: uniform, 64  $\lambda$  per link




(b) After: demand-aware (floor 0.1, cap  $2\times$ )



orange = hot-expert GPUs (load  $\geq 1.5\times$ ); blue = source G0; arrow width  $\propto \lambda$

(c) Source G0:  $\lambda$  per outgoing link



(d) After: active drop rings per GPU

